

HashFilter: Filtering Streams of Data for Top-K Heavy Hitters

Justin Athill, Angela Hillsman, and Daniel Wood

Abstract

In this paper, we present two implementations of HashFilter, a heavy hitter detection algorithm that identifies the k most frequent senders along an ISP backbone link for the purpose of DoS detection. We address the advantages and disadvantages of sampling techniques and counter based methods, and propose an improved algorithm that increases accuracy per space used by combining elements of both. Furthermore, we design HashFilter to run on emerging programmable switches. HashFilter is prototyped in P4 and evaluated using 3 separate CAIDA datasets from an ISP backbone link. For top- k , we experimentally identify over 98 percent of the top 300 hitters using 4500 counters on a trace containing nearly 3 million packets.

1 Introduction

The Heavy Hitter problem refers to the objective of identifying the heaviest flows in a stream of data. In one variant of the problem, heavy flows are classified as those with a frequency above a threshold t . In this paper, we address a second variant of the problem—“top- k .” In this variant, the heavy hitters are the top k flows by frequency.

There are many potential flows that can be analyzed in the context of the Heavy Hitters problem, including source IP addresses, destination IP addresses, transport port numbers, or five-tuples. Depending on the application of the algorithm, a different conception of flow may be appropriate. In this case, we employ our Heavy Hitters algorithm for DoS detection by identifying hosts that are responsible for sending the most traffic through an ISP link. Due to the nature of Internet traffic, a relatively few number of hosts are responsible for sending the majority of packets through a network. According to Figure 1, the top 3 percent of hosts may account for over half of the packets traveling




Figure 1: Graph of the cumulative traffic addressed to/from the top k sources/destinations, captured from an ISP backbone link.

through an ISP link in a given time period. The graph shows the cumulative traffic addressed to/from the top k sources/destinations, and reveals two important conclusions about the distribution of traffic. First, there are approximately twice as many destinations as sources, and as a result, these sources account for more traffic on average. Second, for both source addresses and destination addresses, the heaviest hitters are responsible for a highly disproportionate amount of network traffic. Through the rest of this paper, we consider the frequency of packets identified by source IP address as our measure of heavy hitters.

2 Related Work

2.1 Sampling Algorithms

Rather than count the frequency of each flow, one family of heavy hitter algorithms is based on infrequent sampling. While this approach improves scalability and drastically reduces memory usage, sampling comes at the cost of decreased accuracy. The algorithms Sampled NetFlow [7] and Sample and Hold [4] follow this approach by sampling each packet with some very small probability, such as 0.1 percent or even 0.01 percent. Sample and Hold improves upon Netflow, since once a flow is sampled, a corresponding counter is held in a hash table in

flow memory until the end of the measurement interval. The entry for a flow is updated for every subsequent packet belonging to the flow. According to this scheme, the error is proportional to $1/M$, as opposed to $1/M$ for a classical sampling algorithm, where M is the available memory.

2.2 Sketch Algorithms

Sketch algorithms are a family of algorithms that attempt to answer questions about data streams by leveraging approximation. By sacrificing accuracy, sketch algorithms allow the heavy hitter problem to be tackled using a limited amount of memory. One of the most prominent of these algorithms is Count-Min Sketch [3], which uses a two-dimensional hash table to approximate counts. Each packet is hashed using d different hash functions, and the counter in each of the d corresponding buckets is incremented. Hash collisions will cause certain bucket counts to be incremented for different flow identifiers, so to approximate the count for a given identifier, it is hashed using the d hash functions, and the minimum of these bucket counts is used. This method is hardware-friendly, but sketch based algorithms do not track the flow identifiers associated with each count, so reporting the frequency of the top k flows is not accurate.

2.3 Counter Based Algorithms

Counter Based Algorithms process every packet, but due to memory constraints, are only able to maintain a counter of a constant number of the heaviest flows. Therefore, these algorithms aim to retain counters of only heavy hitters while ignoring lighter flows through the use of strategic admission and eviction policies. Space Saving (Algorithm 1) [6] maintains a table of the frequencies of N flows, evicting the least frequent flow each time an unmonitored flow is encountered. The newly admitted element assumes the frequency of the evicted flow. This eviction policy results in large errors for heavy-tailed workloads, where many new small flows may wrongly evict larger flows. Furthermore, depending on the hardware implementation, it may be resource intensive to find the least frequent flow in the table for every newly encountered flow.

The intuition behind Randomized Admission Policy (Algorithm 2) [1] is to minimize this error by be-

Algorithm 1: Space Saving Algorithm

```

1 Table  $T$  has  $m$  slots, either containing  $(key_j, val_j)$ 
  at slot  $j \in \{1, \dots, m\}$ , or empty. Incoming packet
  has key  $iKey$ 
2 if  $\exists$  slot  $j$  in  $T$  with  $iKey = key_j$  then
3    $val_j \leftarrow val_j + 1$ 
4 else
5   if  $\exists$  empty slot  $j$  in  $T$  then
6      $(key_j, val_j) \leftarrow (iKey, 1)$ 
7   else
8      $r \leftarrow \operatorname{argmin}_{j \in \{1, \dots, m\}} (val_j)$ 
9      $(key_r, val_r) \leftarrow (iKey, val_r + 1)$ 
10  end
11 end

```

ing much more conservative about the elements that are admitted into the table. Instead of evicting the minimum element in the table for every new flow encountered, RAP only admits new flows with a probability of $1/(c_m + 1)$, where c_m is the minimal counter value. By being more conservative, RAP has increased accuracy over Space Saving, but also makes it more difficult for new larger flows to gain admission. In our approach, we adopt elements of the randomized admission policy to prevent false evictions while also dampening the adverse effects of a conservative policy.

HashPipe [8] is heavily inspired by Space Saving, but leverages feed-forward packet processing to divide the task of finding the minimum into small parts. The algorithm consists of d stages, each with its own hash function h_d and an associated hash table. When a packet enters the pipeline at the first stage, the hash function h_0 is used to hash its identifier to a bucket. Each bucket will contain a flow identifier and its associated count. If the identifier of an incoming packet matches the identifier stored in the bucket it hashes to, the count is incremented, and no further processing is done. Similarly, if the bucket is empty, the identifier is added with a count of 1, and processing stops. However, the rest of the pipeline comes into play if the incoming packet identifier does not match the stored packet identifier. In this case, the stored packet will be evicted, and the incoming packet will be stored in its place with a count of 1.

In subsequent stages, the key-counter pair that has been evicted is hashed using the corresponding hash function h_d and it is compared to the stored value in that bucket. Instead of always evicting like in the first stage, the stored value will only be evicted if it is the minimum between the two counter values. If the stored value is less, the key-counter pairs are

Algorithm 2: HashPipe: Pipeline of d hash tables [6]

```

1      > Insert in the first stage
2   $L_1 \leftarrow h_1(iKey)$ 
3  if  $key_{L_1} = iKey$  then
4       $val_{L_1} \leftarrow val_{L_1} + 1$ 
5      end processing
6  end
7  else if  $L_1$  is an empty slot then
8       $(key_{L_1}, val_{L_1}) \leftarrow (iKey, 1)$ 
9      end processing
10 end
11 else
12      $(cKey, cVal_j) \leftarrow (key_{L_1}, val_{L_1})$ 
13      $(key_{L_1}, val_{L_1}) \leftarrow (iKey, 1)$ 
14 end
15     > Track a rolling minimum
16 for  $i \leftarrow 2$  to  $d$  do
17      $L \leftarrow h_i(cKey)$ 
18     if  $key_L = cKey$  then
19          $val_L \leftarrow val_L + cVal$ 
20         end processing
21     end
22     else if  $L$  is an empty slot then
23          $(key_L, val_L) \leftarrow (cKey, cVal)$ 
24         end processing
25     end
26     else if  $val_L < cVal$  then
27         swap  $(cKey, cVal)$  with  $(key_L, val_L)$ 
28     end
29 end

```

swapped, and the key previously in the table is carried to the next stage. This same process is repeated at each stage until one pseudo-minimum value is carried off the end.

Instead of attempting to find a true minimum, as is the case in Space Saving, HashPipe settles for a probabilistic minimum, obtained by comparing only one value per stage. The main idea behind the pipeline is that heavy flows will be retained and lighter flows will be evicted over time. HashPipe is useful because for each packet, there is only one read per table. This allows for efficient stream processing and gets around hardware constraints that do not allow writes based on multiple reads in the same table. The downside of this scheme is that it does not prevent duplicate keys across different tables. When a count is stored for the same keys in different stages, this reduces the space available to hold onto heavier flows. However, duplicates have been shown to account for only 5-10 percent of table space, and have a limited impact on accuracy [8]. With a fixed amount of memory available to the hardware, the number of stages d can be tuned: a greater number of stages increases heavy flow retention because more slots are sampled to pick a minimum, but it will increase the number of duplicates.




Figure 2: The minimum of a randomly selected set of elements approaches the true minimum.

3 Design

3.1 Probabilistic Minimum

Starting with the Space Saving algorithm as a baseline, we first settled on probabilistic sampling as a way to make finding the minimum element more efficient, rather than linearly searching the entire table. As long as a flow with a low enough frequency is evicted when a new flow is encountered, larger flows will still be preserved in the table. As Figure 2 shows, sampling as few as 4 elements probabilistically ensures that only elements in the lowest quintile will be evicted. Drawing from HashPipe, we implemented this optimization using a hash table pipeline, as the number of stages equals the number of elements probabilistically sampled. If there are four stages in the pipeline, the minimum of the four flows encountered will be a candidate for eviction.

3.2 Randomized Assignment

HashPipe, as the name suggests, features a consistent hashing scheme whereby each stage uses an independent hash function and subsequent repeated flows consistently hash to the same location within each stage. This scheme is intended to reduce the number of duplicate flows occupying space in the pipeline and consolidate their counts. However, we experimentally determined that using a random function to determine flow assignment within the first stage resulted in improvements in accuracy of nearly 50 percent.

3.3 Logarithmic Admission Policy

One of the major flaws with HashPipe is the fact that, as with Space Saving, every flow will always be admitted into the pipeline, even if it will never reoccur. In streams with many small flows, accuracy suffers as heavy hitters near the minimum threshold are evicted by small newly encountered flows. Therefore, in HashFilter, we introduced a modified admission policy inspired by RAP, but instead of making the probability of admission inversely proportional to the frequency of the minimum counter, we used a log function to dampen extremely low admission probabilities. We introduced a log function to counteract one of the systematic errors with RAP: denying admission to new heavy hitters. This way, flows with high frequencies are still protected, but smaller flows have a larger probability of being evicted. We tested two potential solutions to this problem, both of which restrict admission to the table. We call these variants of HashFilter Front Rejection and Back Rejection. Our implementation can be accessed at github.com/danielwood95/heavyHitters

In Front Rejection (A.K.A. The Bouncer), packets are randomly assigned a slot in the first stage. If empty, the packet is inserted and the frequency is set to 1. If the slot contains a packet with the identical source IP address, the frequency is incremented. Otherwise, the packet evicts the resident of the slot with probability $p = 1/(5 * \log(c_m + 1))$, where c_m is the frequency count of the flow in the slot. This is easy to implement, and it saves a lot of computational resources further down the pipeline, since on average 90 percent of packets will be refused admission to the pipeline. However, it is not necessarily accurate because the probability of admission is determined by only the flow randomly compared to in the first stage of the pipeline. This flow is not guaranteed to be the minimum, but it is likely to be an average flow, which turns out to be an adequate compromise. The pseudo code is detailed in Algorithm 3.

In Back Rejection (A.K.A. The Interview), all packets are admitted to the entrance of the pipeline and proceed through all stages of the pipeline. However, rather than always evicting the minimum of the flows encountered throughout the pipeline to make way for the new flow, a calculation is made at the

Algorithm 3: HashPipe Bouncer Admission

```

1  > Randomly position in first stage
2   $L_1 \leftarrow \text{rand}_i(\text{iKey})$ 
3  if  $\text{key}_{L_1} = \text{iKey}$  then
4     $\text{val}_{L_1} \leftarrow \text{val}_{L_1} + 1$ 
5    end processing
6  end
7  else if  $L_1$  is an empty slot then
8     $(\text{key}_{L_1}, \text{val}_{L_1}) \leftarrow (\text{iKey}, 1)$ 
9    end processing
10 end
11 else
12   > Randomized Admission Policy
13    $m \leftarrow \text{val}_{L_1}$ 
14   if  $\text{random}() < \frac{1}{5 * \log(m+1)}$  then
15      $(\text{cKey}, \text{cVal}_j) \leftarrow (\text{key}_{L_1}, \text{val}_{L_1})$ 
16      $(\text{key}_{L_1}, \text{val}_{L_1}) \leftarrow (\text{iKey}, 1)$ 
17   end
18   else
19     end processing
20 end
21   > Track a rolling minimum
22 for  $i \leftarrow 2$  to  $d$  do
23    $L \leftarrow h_i(\text{cKey})$ 
24   if  $\text{key}_L = \text{cKey}$  then
25      $\text{val}_L \leftarrow \text{val}_L + \text{cVal}$ 
26     end processing
27   end
28   else if  $L$  is an empty slot then
29      $(\text{key}_L, \text{val}_L) \leftarrow (\text{cKey}, \text{cVal})$ 
30     end processing
31   end
32   else if  $\text{val}_L < \text{cVal}$  then
33     swap  $(\text{cKey}, \text{cVal})$  with  $(\text{key}_L, \text{val}_L)$ 
34   end
35 end

```

end of the pipeline. With probability $p = 1/(5 * \log(c_m + 1))$, the minimum flow is evicted, otherwise the minimum flow is inserted back into the first stage, retroactively denying admission to the newly encountered flow located in the first stage.

4 Prototype in P4

To show that our algorithms are implementable in modern programmable switch languages, we implemented Front Rejection and Back Rejection in P4 version 1.1 [2] by modifying the P4 implementation of HashPipe [9]. The HashPipe implementation of P4 uses a custom hash function with P4's `modify_field_with_hash_based_offset()` to hash each IP source address to an index of a hash table, which is implemented as a stateful register in the switch. The HashPipe implementation also uses metadata to track the carried key and count throughout each stage. We use this setup for each of our stages as well.

To implement Front Rejection, on the first stage we use P4's random number generator, `modify_field_rng_uniform()`, to generate

an index. To implement our placement probability $p = 1/(5 * \log(c_m + 1))$, where c_m is the count at the randomly generated index, we do two things. As the logarithm function cannot be implemented with P4’s basic bitwise functions and inability to loop, we first approximate logarithm into 16 “buckets” with sizes ranging from 1 to nearly 100,000 as the change per step decreases. For example, $c_m = 16 - 20$ ranges from $p = 0.1625 - 0.1513$, and are consolidated into a bucket with value 16. We confirmed that these buckets have nearly identical results by running each of our simulations, discussed in the next section, with the approximate logarithm function. We then generate a second random integer between 0 and 100 and compare it to our approximate logarithm value: if the random integer is greater we keep the current value and otherwise we replace it. In the second and subsequent stages, we use the same implementation as HashPipe. We implemented this with only two stages, as all additional stages would be a duplicate of stage 2.

We created two implementations of Back Rejection. In both implementations’ first stage, we save the hashed index of the original packet in our metadata. In the first version, to avoid doing a second loop, we create an additional stage. This stage revisits the original hashed index in stage 1’s table and replaces it with probability $p = 1/(5 * \log(c_m + 1))$ using the same method discussed above. Though we did not encounter any issues with the second access of stage 1’s table registers, we were uncertain if this would cause issues in some hardware. Thus, in our second version we recirculate the packet. In the final stage of the first loop, we indicate in the metadata that it has completed a first pass and use P4’s `resubmit()` function to place it back on the ingress pipeline. To prevent the packet from traversing through all stages, we take advantage of P4’s `match:action` structure to assure it only goes through the additional stage mentioned in the first version. While recirculation can effectively halve the available packet-processing bandwidth, we believe the impact is less since only one stage is completed twice. This impact could also potentially be avoided altogether with the first version. We implemented both versions with only three stages, as all additional stages would be a duplicate of stage 2.

All three switch implementations have been successfully compiled and brought up on a Mininet [10]

instance, but due to complications with the Mininet connections were unable to run with simulated traffic. While simulation of traffic would confirm proper function of each implementation, it would be impossible to exactly match simulation results due to our inability to perfectly mock P4’s random number generator. These implementations still confirm that our algorithms can be implemented with programmable hardware today.

5 Evaluation

We evaluate the accuracy of HashFilter through a series of simulations in which we fine tune the various parameters: number of stages, memory size, hash functions vs. random assignment, and probabilistic admission coefficients. In order to simulate realistic streams of traffic, we ran testing using three different traces from the equinix chicago ISP backbone link, recorded in 2016. These anonymized traces each contain between 20 - 40 million packets, over 1 million different flows, and range from 40 minutes to 1 hour long. The data was obtained with permission from the Center for Applied Internet Data Analysis (CAIDA) [5]. We parsed the data from the CAIDA traces to isolate only the source IP addresses. Each source IP address is considered a separate flow.

5.1 Accuracy Metrics

Through measuring the false negative rate when testing our algorithms against the CAIDA data, we attempted to experimentally determine which combination of design policies yielded the best results. We focused on testing the two different admission policies, and applied different levels of randomized slot assignment to find an optimal algorithm. When applying front rejection, we found that uniformly randomizing the table index only in the first stage yielded the best results, and gave as much as a 50 percent accuracy improvement over simple front rejection with normal hashing (see Figure 3). While some randomization provided improvements for front rejection, our tests of back rejection showed that normal hashing at all stages was best (see Figure 4). Why is there a discrepancy in the effectiveness of random slot assignment? Perhaps this can be explained by the increased level of eviction that occurs in the first stage when random indices are used.

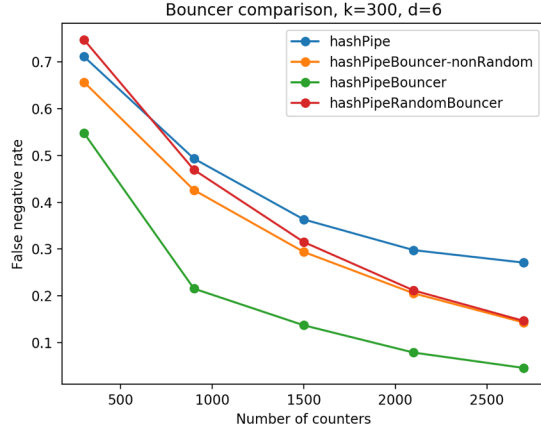


Figure 3: Comparison of Bouncer admission policy algorithm with different combinations of randomization and hashing at each stage. In hashPipeBouncer-nonRandom, consistent hash functions are used in every stage. The hashPipeBouncer algorithm only applies randomness in the first stage. This variation performed best across all memory sizes and is the main Bouncer algorithm we use in further comparison. The hashPipeRandomBouncer algorithm uses randomness in all stages.




Figure 4: Impact of random assignment versus hash functions. The hashPipeRandom algorithm randomizes indices in all stages and provides some benefits at high memory sizes. The hashPipefRand algorithm assigns a random index in the first stage and then uses consistent hash functions in the following stages. Randomization in the first stage improves accuracy rates across the entire tested memory range, showing that randomization is useful absent any strict admission policy.

When slot assignments are randomized, it is less likely that the same key will end up at the same index and increase its count in the first stage. So effectively, keys are more quickly pushed to later stages where there is a greater emphasis on retaining heavy flows. Since consistent hashing is performed in later stages, the same key is able to add to its count, and can do this more frequently when it is quickly sent down the pipeline from the first stage. Randomized slot assignment more effectively treats the first stage as transient, so light flows will more quickly be evicted and heavy flows will more quickly be retained in the core of the pipeline. For this reason, randomizing indices in all stages is less effective because retention also is lower in the latter stages, when it is most important to hold onto heavy flows.

Randomization in the first stage adds an accuracy boost to front rejection, but it introduces problems with the more complex back rejection policy. The goal with the Interview policy, is to deny first stage admission to the incoming key at a high rate, and replace it with the probabilistic minimum that emerges from the pipeline. However, when randomization is applied, we can't guarantee that we are swapping the minimum with the incoming key in the first stage. This renders our stricter admission policy ineffective

if we can't guarantee that the incoming key is the one being evicted in the first stage. Evidently, randomized slot assignment in the first stage can effectively increase accuracy under the right conditions, and was even shown to be useful as a standalone improvement to HashPipe without applying front or back rejection (see Figure 5). However, front rejection does not benefit from randomized slot assignment, and the standalone version of the policy ended up with the best performance. In all further testing when randomization is applied, we run multiple trials on each dataset and report average values. We were also able to tune our algorithm's performance by experimenting with the multiplicative factor in our logarithmic admission equation. Figure 6 shows how performance improved up to a log factor value of 5, and experienced diminishing returns thereafter. We settled on using a log factor of 5 for both front and back rejection admission policies for all further HashFilter testing.

5.2 Comparison of HashPipe Implementations

We began algorithm comparison by varying the available memory size, which corresponds to a




Figure 5: Impact of random slot assignment versus hash functions. The hashPipeRandom algorithm randomizes indices in all stages and provides some benefits at high memory sizes. The hashPipefRand algorithm only hashes to a random index in the first stage and then uses consistent hash functions in the following stages. Randomization in the first stage improves accuracy rates across the entire tested memory range, showing that randomization is useful absent any strict admission policy.

greater number of counters that can be stored across all hash tables. Figure 7 shows our results when searching for the top 300 flows and using 6 table stages. Both algorithms that applied an admission policy outperformed the standard HashPipe algorithm for all tested memory sizes. Accuracy rates improved by about 20 percent when the number of counters was limited to 300, equal to the number of top- k flows being identified, and increases to more than 50 percent improvement when more than 2000 counters are used. HashPipe with front rejection in particular brought its accuracy rate to more than 90 percent with 2000 counters, which is less than half the memory required by the standard HashPipe algorithm to achieve that accuracy with 300 heavy hitters and 6 table stages. In all algorithms, improvement begins to diminish after 1000 counters are available. In addition to testing the impact of varying memory, we also experimented with the number of pipeline stages available. HashPipe tends to perform best with a limited amount of stages so that the number of duplicates is limited. The optimal number of stages is about 6 with the standard HashPipe, but Figure 8 shows that our algorithms continue to improve past this point as the number of table stages is increased. While Interview and Bouncer also experience dimin-

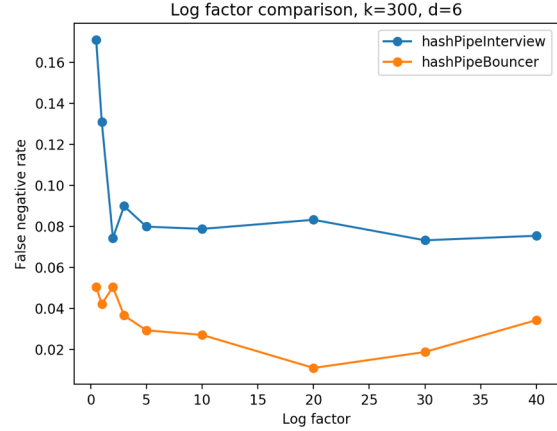


Figure 6: Impact of varying the factor f when calculating the admission threshold $p = 1/(f * \log(c_m + 1))$. Results show that factors of 5 and greater experienced relatively similar accuracy rates for both front rejection and back rejection.




Figure 7: Comparison of false negatives of Interview and Bouncer to baseline. Both algorithm optimizations improve on the standard HashPipe algorithm over the entire tested memory range.

ishing returns after about 8 table stages, they do not see the same performance decrease present in the standard HashPipe algorithm.

5.3 Space Saving Comparison

As a final benchmark, we compare both variants of HashFilter to the Space Saving algorithm, which finds the true minimum of all of its counters. While this algorithm allows for more accuracy, hardware constraints preventing writes based off of multiple reads make implementations of this algorithm unre-




Figure 8: Comparison of algorithms when number of table stages is varied. Interview and Bouncer do not suffer from the same accuracy drawbacks as the baseline when increasing the number of table stages.

alistic. Figure 9 shows that our algorithms fall behind an idealized implementation of Space Saving, where it is no issue to compute a global minimum. Similar to the baseline implementation of HashPipe, our implementations also outperform Space Saving when a low amount of memory is available. This is because Space Saving is only guaranteed to hold onto the k th heaviest item when it is larger than average count in the table, so its full benefits are realized when memory is increased




Figure 9: Comparison of Interview and Bouncer with Space Saving with number of counters varied.

6 Conclusions

We set out to create an algorithm that solves the top- k Heavy Hitters problem while making efficient use

of switch hardware functionality. Building off of HashPipe, we created an algorithm, HashFilter, with two variants of restricted admission policy: Bouncer and Interview. These achieved the same accuracy of HashPipe while requiring less than half of the memory resources. Using real traffic data, we show experimental results that justify our design choices and prototype HashFilter with P4 to confirm its viability on hardware.

References

- [1] E. G. F. R. K. Y. Ben Basat, R. Randomized admission policy for efficient top- k and frequency estimation. In *IEEE INFOCOM 2017 Conference on Computer Communications*, Atlanta, GA, 2017.
- [2] T. P. L. Consortium. The p4 language specification: Version 1.1.0, 2016. URL <https://p4.org/p4-spec/p4-14/v1.1.0/tex/p4.pdf>.
- [3] G. Cormode and S. Muthukrishnan. An improved data stream summary: the count-min sketch and its applications. *Journal of Algorithms*, 55(1):58–75, 4 2003.
- [4] C. Estan and G. Varghese. New directions in traffic measurement and accounting: Focusing on the elephants, ignoring the mice. *ACM Transactions on Computer Systems*, 21(3):270–313, 8 2003.
- [5] C. for Applied Internet Data Analysis. Caida data, 2017. URL www.caida.org.
- [6] A. D. Metwally, A. and A. El Abbadi. Efficient computation of frequent and top- k elements in data streams. *International Conference on Database Theory*, 3363(1), 1 2005.
- [7] C. Netflow. Netflow. URL <https://www.cisco.com/c/en/us/products/ios-nx-os-software/ios-netflow/index.html>.
- [8] N. S. R. O. M. S. Sivaraman, V. and J. Rexford. Heavy-hitter detection entirely in the data plane. In *Proceedings of ACM Symposium on SDN Research 2017*, Santa Clara, CA, Apr. 2017.
- [9] V. Sivaraman. P4 implementation of hashpipe code, 2016. URL https://github.com/vibhaa/iw15-heavyhitters/p4v1_1.
- [10] M. Team. Mininet: An instant virtual network on your laptop (or other pc). URL <https://mininet.org>.